

```
In [1]: # This notebook prepares the dataframes for analysis and runs  
# the Random Forest on the combined model (economic and sentiment features) and  
# the economic (only) model.
```

```
In [2]: import warnings  
warnings.filterwarnings('ignore')
```

```
In [3]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn.ensemble import RandomForestRegressor  
from sklearn.model_selection import GridSearchCV  
from sklearn.model_selection import TimeSeriesSplit  
from sklearn.metrics import r2_score, mean_absolute_error
```

```
In [4]: # Quick look at the first two lines of the econ CSV  
with open('data/econ_vars.csv') as f:  
    print(f.readline())  
    next(f)  
    print(f.readline())
```

```
month_date,ur,pi_gr,nyse_gr,gdp_log_diff,tb_sqrt_diff,ics_log_diff,pce_gr_diff
```

```
2011-01-01,9.1,0.002028679786119758,0.016528644263536174,-  
0.0024113163305852225,0.013132595943347536,-0.004034975212179326,-  
0.004803188602163777
```

```
In [5]: # Set data types for the econ features  
dts = {"month_date": str, "ur": np.float64  
      , "pi_gr": np.float64, "nyse_gr": np.float64  
      , "gdp_log_diff": np.float64, "tb_sqrt_diff": np.float64  
      , "ics_log_diff": np.float64, "pce_gr_diff": np.float64}
```

```
In [6]: # Import econ data  
data_raw = pd.read_csv("data/econ_vars.csv"  
                      , sep=","  
                      , skiprows=0  
                      , dtype=dts)
```

```
In [7]: # Check the shape of the raw econ data  
data_raw.shape
```

```
Out[7]: (91, 8)
```

```
In [8]: # Copy raw data  
data = data_raw.copy()
```

```
In [9]: # Sort and reset index  
data.sort_index(ascending=False, inplace=True)  
data.reset_index(drop=True, inplace=True)
```

```
In [10]: # Check the dataframe
data.head()
```

Out[10]:

	month_date	ur	pi_gr	nyse_gr	gdp_log_diff	tb_sqrt_diff	ics_log_diff	pce_gr_diff
0	2018-06-01	4.0	0.003863	-0.001827	0.000000	0.014587	0.002039	-0.001110
1	2018-05-01	3.8	0.003728	0.000941	0.010188	0.037168	-0.008130	0.008168
2	2018-04-01	3.9	-0.004827	-0.002903	0.010188	0.022809	-0.025975	-0.008891
3	2018-03-01	4.1	0.003938	-0.015846	0.000000	0.050844	0.016907	0.007407
4	2018-02-01	4.1	0.003282	-0.053517	0.005479	0.065562	0.040947	0.001663

```
In [11]: # Quick look at the first two lines of the sentiment CSV
with open('data/rev_means_vars_stationary.csv') as f:
    print(f.readline())
    next(f)
    print(f.readline())
```

date,perc_pos_rev_weighted,perc_neg_rev_weighted,perc_uncert_rev_weighted,perc_litig_rev_weighted,perc_modal_wk_rev_weighted,perc_modal_mod_rev_weighted,perc_constraint_rev_weighted,perc_modal_str_rev_weighted_diff

2011-03-01,1.9607048390000001,0.979452392,0.957424913,0.073102265,0.387478896,0.560045154,0.142131806,-0.0495581759999999

```
In [12]: # Import sentiment data
sent_raw = pd.read_csv("data/rev_means_vars_stationary.csv"
    , sep=","
    , skiprows=0
    #, dtype=dfs
    , usecols=[0,1,2,3,4,5,6,7,8]
    )
```

```
In [13]: # Check the sentiment data types
sent_raw.dtypes
```

```
Out[13]: date                                object
perc_pos_rev_weighted                       float64
perc_neg_rev_weighted                       float64
perc_uncert_rev_weighted                   float64
perc_litig_rev_weighted                    float64
perc_modal_wk_rev_weighted                 float64
perc_modal_mod_rev_weighted                float64
perc_constraint_rev_weighted               float64
perc_modal_str_rev_weighted_diff           float64
dtype: object
```

```
In [14]: # Drop any observations with Nan
sent_raw.dropna(inplace=True)
```

```
In [15]: #sort and reset index
sent_raw.sort_values(by='date', ascending=False, inplace=True)
sent_raw.reset_index(drop=True, inplace=True)
```

```
In [16]: # Confirm the shape of the dataframe
sent_raw.shape
```

```
Out[16]: (89, 9)
```

```
In [17]: # Check the dataframe
sent_raw.head()
```

```
Out[17]:
```

	date	perc_pos_rev_weighted	perc_neg_rev_weighted	perc_uncert_rev_weighted	perc_litig
0	2018-06-01	2.043836	0.890239	0.815790	0.094844
1	2018-05-01	1.915107	1.074677	0.803269	0.082565
2	2018-04-01	1.926137	1.068873	0.777739	0.095691
3	2018-03-01	2.016781	0.979256	0.800769	0.102810
4	2018-02-01	2.004129	0.963727	0.815629	0.098500

```
In [18]: # Copy raw data
sent = sent_raw.copy()
```

```
In [19]: # Filter econ observations to match the sentiment dataframe
data = data.iloc[0:89,:]
```

```
In [20]: # Combine the econ and sentiment dataframes
data = pd.concat([data,sent], axis=1)
```

```
In [21]: # Replace the integer index with the date column
data = data.set_index("date")
```

```
In [22]: # Create response variable
y_true = data['pce_gr_diff']
```

```
In [23]: # Drop the month_date column and the gdp column
data = data.drop(data.columns[[0,4,7]], axis=1)
```

```
In [24]: # Check the dataframe
data.head()
```

```
Out[24]:
```

	ur	pi_gr	nyse_gr	tb_sqrt_diff	ics_log_diff	perc_pos_rev_weighted	perc_neg_rev
date							
2018-06-01	4.0	0.003863	-0.001827	0.014587	0.002039	2.043836	0.890239
2018-05-01	3.8	0.003728	0.000941	0.037168	-0.008130	1.915107	1.074677
2018-04-01	3.9	-0.004827	-0.002903	0.022809	-0.025975	1.926137	1.068873
2018-03-01	4.1	0.003938	-0.015846	0.050844	0.016907	2.016781	0.979256
2018-02-01	4.1	0.003282	-0.053517	0.065562	0.040947	2.004129	0.963727

```
In [25]: # Copy econ features to econ dataframe and drop all sentiment variables
econ = data.drop(data.columns[[5,6,7,8,9,10,11,12]], axis=1)
```

```
In [26]: # Check the dataframe
econ.head()
```

Out[26]:

	ur	pi_gr	nyse_gr	tb_sqrt_diff	ics_log_diff
date					
2018-06-01	4.0	0.003863	-0.001827	0.014587	0.002039
2018-05-01	3.8	0.003728	0.000941	0.037168	-0.008130
2018-04-01	3.9	-0.004827	-0.002903	0.022809	-0.025975
2018-03-01	4.1	0.003938	-0.015846	0.050844	0.016907
2018-02-01	4.1	0.003282	-0.053517	0.065562	0.040947

```
In [27]: # create time series cross validation split
tscv = TimeSeriesSplit(n_splits=3)
print(tscv)

TimeSeriesSplit(max_train_size=None, n_splits=3)
```

```
In [28]: # instantiate random forest object
rfreg = RandomForestRegressor()
```

```
In [29]: # establish parameters for grid search object
parameters = {"n_estimators": [50, 100],
              "max_features": [3, 4, 6],
              "min_samples_leaf": [3, 5, 10],
              "max_depth": [3, None],
              "min_samples_split": [2, 5, 10],
              "bootstrap": [True],
              "random_state": [987]
             }
```

```
In [30]: # create a grid search object
comb_model_gs = GridSearchCV(estimator=rfreg
                             , n_jobs=-1 # parallel execution -1 is all processors
                             , verbose=1 # low verbosity
                             , param_grid=parameters
                             , cv=tscv # KFold = 3
                             , scoring="neg_mean_squared_error"
                             )
```

```
In [31]: # fit random forest regressor
comb_model_gs.fit(data, y_true)
```

Fitting 3 folds for each of 108 candidates, totalling 324 fits

```
[Parallel(n_jobs=-1)]: Done 32 tasks      | elapsed:    0.8s
[Parallel(n_jobs=-1)]: Done 324 out of 324 | elapsed:    4.3s finished
```

```
Out[31]: GridSearchCV(cv=TimeSeriesSplit(max_train_size=None, n_splits=3),
                    error_score='raise',
                    estimator=RandomForestRegressor(bootstrap=True, criterion='mse', max_dept
h=None,
                    max_features='auto', max_leaf_nodes=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=1, min_samples_split=2,
                    min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                    oob_score=False, random_state=None, verbose=0, warm_start=False),
                    fit_params=None, iid=True, n_jobs=-1,
                    param_grid={'n_estimators': [50, 100], 'max_features': [3, 4, 6], 'min_sa
mples_leaf': [3, 5, 10], 'max_depth': [3, None], 'min_samples_split': [2, 5,
10], 'bootstrap': [True], 'random_state': [987]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                    scoring='neg_mean_squared_error', verbose=1)
```

```
In [32]: # print grid search results sorted by mean test score
comb_model_results = pd.DataFrame(comb_model_gs.cv_results_)
headers = ['rank_test_score', 'mean_test_score', 'mean_train_score', 'mean_fit_time',
            'split0_train_score', 'split1_train_score', 'split2_train_score']
comb_model_results.sort_values(by=["rank_test_score"], inplace=True, ascending=True)
comb_model_results[headers].head(10)
```

Out[32]:

	rank_test_score	mean_test_score	mean_train_score	mean_fit_time	split0_train_score	split1_train_score
93	1	-0.00002	-0.000008	0.194232	-0.000011	-0.000007
91	1	-0.00002	-0.000008	0.203017	-0.000011	-0.000007
39	3	-0.00002	-0.000010	0.182110	-0.000011	-0.000009
37	3	-0.00002	-0.000010	0.176719	-0.000011	-0.000009
43	5	-0.00002	-0.000015	0.177317	-0.000021	-0.000012
47	5	-0.00002	-0.000015	0.183128	-0.000021	-0.000012
45	5	-0.00002	-0.000015	0.175593	-0.000021	-0.000012
42	8	-0.00002	-0.000015	0.092832	-0.000020	-0.000012
44	8	-0.00002	-0.000015	0.087386	-0.000020	-0.000012
46	8	-0.00002	-0.000015	0.093360	-0.000020	-0.000012

```
In [33]: # View the optimal random forest model parameters
comb_model_gs.best_estimator_
```

```
Out[33]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                               max_features=6, max_leaf_nodes=None, min_impurity_decrease=0.0,
                               min_impurity_split=None, min_samples_leaf=3,
                               min_samples_split=2, min_weight_fraction_leaf=0.0,
                               n_estimators=100, n_jobs=1, oob_score=False, random_state=987,
                               verbose=0, warm_start=False)
```


```
In [34]: print("Combined model MSE: {}".format(comb_model_gs.best_score_))
```

Combined model MSE: -2.0119435819526655e-05

```
In [35]: # Run a new model using the best model from the grid search
comb_best_model = comb_model_gs.best_params_
rfreg_features = RandomForestRegressor(**comb_best_model)
rfreg_features.fit(data, y_true)
```

```
Out[35]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                               max_features=6, max_leaf_nodes=None, min_impurity_decrease=0.0,
                               min_impurity_split=None, min_samples_leaf=3,
                               min_samples_split=2, min_weight_fraction_leaf=0.0,
                               n_estimators=100, n_jobs=1, oob_score=False, random_state=987,
                               verbose=0, warm_start=False)
```

```
In [36]: # Plot the feature importances
feature_importance = rfreg_features.feature_importances_
feature_importance = 100.0 * (feature_importance / feature_importance.max())
sorted_idx = np.argsort(feature_importance)
pos = np.arange(sorted_idx.shape[0]) + .5
plt.figure(figsize=(10, 10))
plt.barh(pos, feature_importance[sorted_idx], align='center', color='navy')
plt.yticks(pos, np.asarray(data.columns.tolist())[sorted_idx]) # data is X
plt.xlabel('Relative Importance')
plt.title('Combined Model Feature Importance')
plt.show()
#plt.savefig('Random Forest Feature Importance.png')
```



```
In [37]: # establish parameters for grid search object
parameters = {"n_estimators": [50, 100],
              "max_features": [2, 3, 4],
              "min_samples_leaf": [3, 5, 10],
              "max_depth": [3, None],
              "min_samples_split": [2, 5, 10],
              "bootstrap": [True],
              "random_state": [987]
            }
```

```
In [38]: # create a grid search object
econ_model_gs = GridSearchCV(estimator=rftree
                             , n_jobs=-1 # parallel execution -1 is all processors
                             , verbose=1 # low verbosity
                             , param_grid=parameters
                             , cv=tscv # KFold = 3
                             , scoring="neg_mean_squared_error"
                             )
```

```
In [39]: # fit random forest regressor
econ_model_gs.fit(econ, y_true)
```

Fitting 3 folds for each of 108 candidates, totalling 324 fits

```
[Parallel(n_jobs=-1)]: Done 32 tasks      | elapsed:    0.8s
[Parallel(n_jobs=-1)]: Done 324 out of 324 | elapsed:    4.2s finished
```

```
Out[39]: GridSearchCV(cv=TimeSeriesSplit(max_train_size=None, n_splits=3),
                    error_score='raise',
                    estimator=RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                    max_features='auto', max_leaf_nodes=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=1, min_samples_split=2,
                    min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                    oob_score=False, random_state=None, verbose=0, warm_start=False),
                    fit_params=None, iid=True, n_jobs=-1,
                    param_grid={'n_estimators': [50, 100], 'max_features': [2, 3, 4], 'min_samples_leaf': [3, 5, 10], 'max_depth': [3, None], 'min_samples_split': [2, 5, 10], 'bootstrap': [True], 'random_state': [987]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                    scoring='neg_mean_squared_error', verbose=1)
```

```
In [40]: # print results sorted by mean test score
econ_model_results = pd.DataFrame(econ_model_gs.cv_results_)
headers = ['rank_test_score', 'mean_test_score', 'mean_train_score', 'mean_fit_time',
           'split0_train_score', 'split1_train_score', 'split2_train_score']
econ_model_results.sort_values(by=["rank_test_score"], inplace=True, ascending=True)
econ_model_results[headers].head(10)
```

Out[40]:

	rank_test_score	mean_test_score	mean_train_score	mean_fit_time	split0_train_score	split1_train_score
53	1	-0.000021	-0.000023	0.172886	-0.000033	-0.000033
105	1	-0.000021	-0.000023	0.134153	-0.000033	-0.000033
103	1	-0.000021	-0.000023	0.141396	-0.000033	-0.000033
51	1	-0.000021	-0.000023	0.174077	-0.000033	-0.000033
49	1	-0.000021	-0.000023	0.181301	-0.000033	-0.000033
107	1	-0.000021	-0.000023	0.104362	-0.000033	-0.000033
50	7	-0.000021	-0.000023	0.089039	-0.000033	-0.000033
52	7	-0.000021	-0.000023	0.093082	-0.000033	-0.000033
106	7	-0.000021	-0.000023	0.056159	-0.000033	-0.000033
48	7	-0.000021	-0.000023	0.092212	-0.000033	-0.000033

```
In [41]: econ_model_gs.best_estimator_
```

```
Out[41]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=3,
                               max_features=4, max_leaf_nodes=None, min_impurity_decrease=0.0,
                               min_impurity_split=None, min_samples_leaf=10,
                               min_samples_split=2, min_weight_fraction_leaf=0.0,
                               n_estimators=100, n_jobs=1, oob_score=False, random_state=987,
                               verbose=0, warm_start=False)
```

```
In [42]: print("Economic model MSE: {}".format(econ_model_gs.best_score_))
```

Economic model MSE: -2.0665388481727326e-05

```
In [43]: # Run a new model using the best model from the grid search
econ_best_model = econ_model_gs.best_params_
rfreg_features = RandomForestRegressor(**econ_best_model)
rfreg_features.fit(econ, y_true)
```

```
Out[43]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=3,
                               max_features=4, max_leaf_nodes=None, min_impurity_decrease=0.0,
                               min_impurity_split=None, min_samples_leaf=10,
                               min_samples_split=2, min_weight_fraction_leaf=0.0,
                               n_estimators=100, n_jobs=1, oob_score=False, random_state=987,
                               verbose=0, warm_start=False)
```



```
In [44]: # Plot the feature importances
feature_importance = rfreg_features.feature_importances_
feature_importance = 100.0 * (feature_importance / feature_importance.max())
sorted_idx = np.argsort(feature_importance)
pos = np.arange(sorted_idx.shape[0]) + .5
plt.figure(figsize=(10, 10))
plt.barh(pos, feature_importance[sorted_idx], align='center', color='maroon')
plt.yticks(pos, np.asarray(econ.columns.tolist())[sorted_idx]) # econ is X
plt.xlabel('Relative Importance')
plt.title('Economic Model Feature Importance')
plt.show()
#plt.savefig('Random Forest Feature Importance.png')
```

