

```
In [1]: # This notebook prepares the dataframes for analysis and runs  
# the VAR on the combined model (economic and sentiment features) and  
# the economic (only) model.
```

```
In [2]: import warnings  
warnings.filterwarnings('ignore')
```

```
In [3]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from statsmodels.tsa.vector_ar.var_model import VAR  
from sklearn.metrics import r2_score  
from statsmodels.stats.stattools import durbin_watson  
from scipy.stats.distributions import chi2  
import warnings  
warnings.filterwarnings('ignore')
```

```
In [4]: import warnings  
warnings.filterwarnings('ignore')
```

```
In [5]: # Quick look at the first two lines of the econ CSV  
with open('data/econ_vars.csv') as f:  
    print(f.readline())  
    next(f)  
    print(f.readline())  
  
month_date,ur,pi_gr,nyse_gr,gdp_log_diff,tb_sqrt_diff,ics_log_diff,pce_gr_diff  
  
2011-01-01,9.1,0.002028679786119758,0.016528644263536174,-  
0.0024113163305852225,0.013132595943347536,-0.004034975212179326,-  
0.004803188602163777
```

```
In [6]: # Set data types for the econ features  
dts = {"month_date": str, "ur": np.float64  
      , "pi_gr": np.float64, "nyse_gr": np.float64  
      , "gdp_log_diff": np.float64, "tb_sqrt_diff": np.float64  
      , "ics_log_diff": np.float64, "pce_gr_diff": np.float64}
```

```
In [7]: # Import econ data  
data_raw = pd.read_csv("data/econ_vars.csv"  
                      , sep=","  
                      , skiprows=0  
                      , dtype=dts)
```

```
In [8]: # Check the shape of the raw econ data  
data_raw.shape
```

```
Out[8]: (91, 8)
```

```
In [9]: # Copy raw data  
data = data_raw.copy()
```

```
In [10]: # Sort descending and reset index  
data.sort_index(ascending=False, inplace=True)  
data.reset_index(drop=True, inplace=True)
```

```
In [11]: # Check the dataframe
data.head()
```

Out[11]:

	month_date	ur	pi_gr	nyse_gr	gdp_log_diff	tb_sqrt_diff	ics_log_diff	pce_gr_diff
0	2018-06-01	4.0	0.003863	-0.001827	0.000000	0.014587	0.002039	-0.001110
1	2018-05-01	3.8	0.003728	0.000941	0.010188	0.037168	-0.008130	0.008168
2	2018-04-01	3.9	-0.004827	-0.002903	0.010188	0.022809	-0.025975	-0.008891
3	2018-03-01	4.1	0.003938	-0.015846	0.000000	0.050844	0.016907	0.007407
4	2018-02-01	4.1	0.003282	-0.053517	0.005479	0.065562	0.040947	0.001663

```
In [12]: # Quick look at the first two lines of the sentiment CSV
with open('data/rev_means_vars_stationary.csv') as f:
    print(f.readline())
    next(f)
    print(f.readline())
```

date,perc\_pos\_rev\_weighted,perc\_neg\_rev\_weighted,perc\_uncert\_rev\_weighted,perc\_l  
itig\_rev\_weighted,perc\_modal\_wk\_rev\_weighted,perc\_modal\_mod\_rev\_weighted,perc\_co  
nstrain\_rev\_weighted,perc\_modal\_str\_rev\_weighted\_diff

2011-03-01,1.9607048390000001,0.979452392,0.957424913,0.073102265,0.387478896,0.  
560045154,0.142131806,-0.0495581759999999

```
In [13]: # Import sentiment data
sent_raw = pd.read_csv("data/rev_means_vars_stationary.csv"
    , sep=","
    , skiprows=0
    #, dtype=dfs
    , usecols=[0,1,2,3,4,5,6,7,8]
    )
```

```
In [14]: # Check the sentiment data types
sent_raw.dtypes
```

```
Out[14]: date                                object
perc_pos_rev_weighted                       float64
perc_neg_rev_weighted                       float64
perc_uncert_rev_weighted                    float64
perc_litig_rev_weighted                     float64
perc_modal_wk_rev_weighted                   float64
perc_modal_mod_rev_weighted                  float64
perc_constrain_rev_weighted                  float64
perc_modal_str_rev_weighted_diff             float64
dtype: object
```

```
In [15]: # Drop any observations with Nan
sent_raw.dropna(inplace=True)
```

```
In [16]: #sort and reset the index
sent_raw.sort_values(by='date', ascending=False, inplace=True)
sent_raw.reset_index(drop=True, inplace=True)
```

```
In [17]: # Confirm the shape of the dataframe
sent_raw.shape
```

```
Out[17]: (89, 9)
```

```
In [18]: # Check the dataframe
sent_raw.head()
```

```
Out[18]:
```

	date	perc_pos_rev_weighted	perc_neg_rev_weighted	perc_uncert_rev_weighted	perc_litig
0	2018-06-01	2.043836	0.890239	0.815790	0.094844
1	2018-05-01	1.915107	1.074677	0.803269	0.082565
2	2018-04-01	1.926137	1.068873	0.777739	0.095691
3	2018-03-01	2.016781	0.979256	0.800769	0.102810
4	2018-02-01	2.004129	0.963727	0.815629	0.098500

```
In [19]: # Copy sentiment data
sent = sent_raw.copy()
```

```
In [20]: # Filter econ observations to match the sentiment dataframe
data = data.iloc[0:89,:]
```

```
In [21]: # Combine the econ and sentiment dataframes
data = pd.concat([data,sent], axis=1)
```

```
In [22]: # Replace the integer index with the date column
data = data.set_index("date")
```

```
In [23]: # Drop the month_date column and the gdp column
data = data.drop(data.columns[[0,4]], axis=1)
```

```
In [24]: # Check the dataframe
data.head()
```

```
Out[24]:
```

	ur	pi_gr	nyse_gr	tb_sqrt_diff	ics_log_diff	pce_gr_diff	perc_pos_rev_weighted	perc_neg_rev_weighted
date								
2018-06-01	4.0	0.003863	-0.001827	0.014587	0.002039	-0.001110	2.043836	0.890239
2018-05-01	3.8	0.003728	0.000941	0.037168	-0.008130	0.008168	1.915107	1.074677
2018-04-01	3.9	-0.004827	-0.002903	0.022809	-0.025975	-0.008891	1.926137	1.068873
2018-03-01	4.1	0.003938	-0.015846	0.050844	0.016907	0.007407	2.016781	0.979256
2018-02-01	4.1	0.003282	-0.053517	0.065562	0.040947	0.001663	2.004129	0.963727

```
In [25]: # Copy econ features to econ dataframe and drop all sentiment variables
econ = data.drop(data.columns[[6,7,8,9,10,11,12,13]], axis=1)
```

```
In [26]: # Check the dataframe
econ.head()
```

Out[26]:

	ur	pi_gr	nyse_gr	tb_sqrt_diff	ics_log_diff	pce_gr_diff
date						
2018-06-01	4.0	0.003863	-0.001827	0.014587	0.002039	-0.001110
2018-05-01	3.8	0.003728	0.000941	0.037168	-0.008130	0.008168
2018-04-01	3.9	-0.004827	-0.002903	0.022809	-0.025975	-0.008891
2018-03-01	4.1	0.003938	-0.015846	0.050844	0.016907	0.007407
2018-02-01	4.1	0.003282	-0.053517	0.065562	0.040947	0.001663

```
In [27]: # Instantiate the VAR model using the combined econ and sentiment variables
comb_model = VAR(data)
```

```
In [28]: # Fit the model to the combined variables
comb_fitted = comb_model.fit(maxlags=2, ic='bic', verbose=True, trend='c')
```

```

VAR Order Selection
=====
          aic          bic          fpe          hqic
-----
0      -91.97      -91.57      1.141e-40      -91.81
1      -99.89*      -93.94*      4.360e-44*      -97.49*
2      -98.88      -87.37      1.649e-43      -94.25
=====
* Minimum

Using 1 based on bic criterion
```

```
In [29]: # Print a summary of the models  
comb_fitted.summary()
```

Out[29]: Summary of Regression Results

```

=====
Model:                                VAR
Method:                               OLS
Date:                                Tue, 18, Jun, 2019
Time:                                17:21:51
-----
No. of Equations:                     14.0000    BIC:                                -93.7304
Nobs:                                88.0000    HQIC:                               -97.2605
Log likelihood:                       2846.13    FPE:                                5.57649e-44
AIC:                                 -99.6422    Det(Omega_mle):                     6.15740e-45
-----
Results for equation ur
=====
=====

```

t-stat	prob	coefficient	std. error
-----			
-----			
const		-0.827841	0.471726
-1.755	0.083		
L1.ur		1.011513	0.011859
85.292	0.000		
L1.pi_gr		1.511155	1.787599
0.845	0.401		
L1.nyse_gr		0.130882	0.435033
0.301	0.764		
L1.tb_sqrt_diff		0.289768	0.323562
0.896	0.373		
L1.ics_log_diff		1.056703	0.304881
3.466	0.001		
L1.pce_gr_diff		-1.561460	2.838331
-0.550	0.584		
L1.perc_pos_rev_weighted		0.085887	0.103277
0.832	0.408		
L1.perc_neg_rev_weighted		0.197955	0.232477
0.852	0.397		
L1.perc_uncert_rev_weighted		0.367334	0.491891
0.747	0.458		
L1.perc_litig_rev_weighted		0.080162	0.983426
0.082	0.935		
L1.perc_modal_wk_rev_weighted		-0.371585	0.736801
-0.504	0.616		
L1.perc_modal_mod_rev_weighted		0.439772	0.461822
0.952	0.344		
L1.perc_constrain_rev_weighted		-0.172817	0.992679
-0.174	0.862		
L1.perc_modal_str_rev_weighted_diff		-0.557680	0.349682
-1.595	0.115		
=====			
=====			

```

Results for equation pi_gr
=====
=====

```

t-stat	prob	coefficient	std. error
-----			
-----			
const		0.023636	0.029573
0.799	0.427		
L1.ur		-0.000248	0.000743
-0.333	0.740		
L1.pi_gr		-0.191584	0.112067
-1.710	0.092		

```
In [30]: # Calculate the combined r squared
y_true = data['pce_gr_diff'].iloc[1:89]
y_pred = comb_fitted.resid['pce_gr_diff'] + y_true
print ("Combined model R^2: {}".format(r2_score(y_true, y_pred)))
```

Combined model R^2: 0.3033409524231241

```
In [31]: # Calculate the mean squared error of the combined model
comb_mse = np.mean(comb_fitted.resid['pce_gr_diff']**2)
print ("Combined model MSE: {}".format(comb_mse))
```

Combined model MSE: 1.746804553927309e-05

```
In [32]: # Run the Durbin Watson test
print ("Combined model Durbin-Watson test: {}".format(durbin_watson(comb_fitted.resid['pce_gr_diff'])))
```

Combined model Durbin-Watson test: 2.5328634408873234

```
In [33]: # Instantiate the VAR model using the econ variables
econ_model = VAR(econ)
```

```
In [34]: # Fit the model to the econ data
econ_fitted = econ_model.fit(maxlags=2, ic='bic', verbose=True, trend='c')
```

#### VAR Order Selection

	aic	bic	fpe	hqic
0	-38.34	-38.17	2.228e-17	-38.27
1	-43.66	-42.47*	1.100e-19	-43.18
2	-44.13*	-41.92	6.936e-20*	-43.24*

\* Minimum

Using 1 based on bic criterion

```
In [35]: # Print a summary of the models  
econ_fitted.summary()
```



Out[35]: Summary of Regression Results

```

=====
Model:                                VAR
Method:                               OLS
Date:                                Tue, 18, Jun, 2019
Time:                                17:21:51
-----
No. of Equations:                     6.00000    BIC:                                -42.4596
Nobs:                                 88.0000    HQIC:                               -43.1656
Log likelihood:                       1213.05    FPE:                                1.11537e-19
AIC:                                  -43.6420    Det(Omega_mle):                     7.04649e-20
-----
Results for equation ur
=====
==

```

	coefficient	std. error	t-stat	pr
ob				
--				
const	0.009825	0.056289	0.175	
0.862				
L1.ur	1.006361	0.008557	117.611	
0.000				
L1.pi_gr	1.210979	1.696079	0.714	
0.477				
L1.nyse_gr	0.131370	0.409557	0.321	
0.749				
L1.tb_sqrt_diff	0.201771	0.293805	0.687	
0.494				
L1.ics_log_diff	0.960370	0.287296	3.343	
0.001				
L1.pce_gr_diff	-1.438920	2.749622	-0.523	
0.602				

```

=====
==
Results for equation pi_gr
=====
==

```

	coefficient	std. error	t-stat	pr
ob				
--				
const	0.004612	0.003593	1.284	
0.203				
L1.ur	-0.000387	0.000546	-0.708	
0.481				
L1.pi_gr	-0.138910	0.108260	-1.283	
0.203				
L1.nyse_gr	0.009626	0.026142	0.368	
0.714				
L1.tb_sqrt_diff	-0.012863	0.018753	-0.686	
0.495				
L1.ics_log_diff	-0.022995	0.018338	-1.254	
0.213				
L1.pce_gr_diff	-0.257348	0.175507	-1.466	
0.146				

```

=====
==
Results for equation nyse_gr
=====
==

```

	coefficient	std. error	t-stat	pr
ob				

```
In [36]: # Calculate the combined r squared
y_true = econ['pce_gr_diff'].iloc[1:89]
y_pred = econ_fitted.resid['pce_gr_diff'] + y_true
print ("Economic model R^2: {}".format(r2_score(y_true, y_pred)))
```

Economic model R^2: 0.28749324079974536

```
In [37]: # Calculate the mean squared error of the econ model
econ_mse = np.mean(econ_fitted.resid['pce_gr_diff']**2)
print ("Economic model MSE: {}".format(econ_mse))
```

Economic model MSE: 1.7865411437689707e-05

```
In [38]: # Run the Durbin Watson test
print ("Economic model Durbin-Watson test: {}".format (durbin_watson(econ_fitted.r
esid['pce_gr_diff'])))
```

Economic model Durbin-Watson test: 2.5575996229732354

```
In [39]: # View the log likelihood values
print ("Combined model Log Likelihood: {}".format(comb_fitted.llf))
print ("Economic model Log Likelihood: {}".format(econ_fitted.llf))
```

Combined model Log Likelihood: 2846.125874607882  
Economic model Log Likelihood: 1213.0463632573587

```
In [40]: # Create a function that calculates the likelihood ratio test
def likelihood_ratio(ll_model_1, ll_model_2): # Model 1 is the more restrictive e
con model
    return (2*(ll_model_2 - ll_model_1))
```

```
In [41]: import math
# Calculate the likelihood ratio and run the likelihood ratio test on a chi square
distribution
LR = likelihood_ratio(econ_fitted.llf, comb_fitted.llf)
pval = chi2.sf(LR, 8) # Combined model +8 variables 8 d.f.
print ("Log Likelihood Ratio test p-value: {:4.03f}".format((pval)))
```

Log Likelihood Ratio test p-value: 0.000